

# INTERPRETABILITY OF BLACK-BOX MODELS

CONTEXT, FORMALIZATION, AND A FOCUS ON COOPERATIVE GAME THEORY

---

<sup>1</sup>EDF Lab Chatou - Département PRISME

<sup>2</sup>Institut de Mathématiques de Toulouse

<sup>3</sup>SINCLAIR AI Lab

*Vendredis de la Data Science*

**MAIF** - Niort.

*2<sup>nd</sup> of December, 2022*

*“ML interpretability methods to certify A.I models on critical systems.”*

*“ML interpretability methods to certify A.I models on critical systems.”*

- **ML interpretability:** Tools that allow to *better understand* a (black-box) model.
- **Critical systems:** Hydroelectric dams, wind turbines, stocks, a human body...
- **A.I. model certification:** Guarantee a certain level of confidence on the A.I via (usually) a set of *meaningful (rigid) rules*.

*“ML interpretability methods to certify A.I models on critical systems.”*

- **ML interpretability:** Tools that allow to *better understand* a (black-box) model.
- **Critical systems:** Hydroelectric dams, wind turbines, stocks, a human body...
- **A.I. model certification:** Guarantee a certain level of confidence on the A.I via (usually) a set of *meaningful (rigid) rules*.

**Main goal:** Develop *meaningful* tools in order to *better understand* a black-box model, in order to *ensure confidence* in their use for *critical systems*.

*“ML interpretability methods to certify A.I models on critical systems.”*

- **ML interpretability:** Tools that allow to *better understand* a (black-box) model.
- **Critical systems:** Hydroelectric dams, wind turbines, stocks, a human body...
- **A.I. model certification:** Guarantee a certain level of confidence on the A.I via (usually) a set of *meaningful (rigid) rules*.

**Main goal:** Develop *meaningful* tools in order to *better understand* a black-box model, in order to *ensure confidence* in their use for *critical systems*.

**Who needs convincing:** *Safety authorities*, which are not receptive to *empirical results* (i.e., SOTA statements).

*“ML interpretability methods to certify A.I models on critical systems.”*

- **ML interpretability:** Tools that allow to *better understand* a (black-box) model.
- **Critical systems:** Hydroelectric dams, wind turbines, stocks, a human body...
- **A.I. model certification:** Guarantee a certain level of confidence on the A.I via (usually) a set of *meaningful (rigid) rules*.

**Main goal:** Develop *meaningful* tools in order to *better understand* a black-box model, in order to *ensure confidence* in their use for *critical systems*.

**Who needs convincing:** *Safety authorities*, which are not receptive to *empirical results* (i.e., SOTA statements).

**Meaningful tools to safety authorities = Theoretical guarantees on the interpretability method.**

1. Explanation vs. Interpretation
2. Interpretability methods
3. Interpretability and cooperative game theory

# The need of a formal definition

From the litterature (Barredo Arrieta et al. 2020):

- **Interpretability:** “[...] *The ability to explain or to provide the meaning in understandable terms to a human.*” - **Provide meaningful (mathematical) tools.**
- **Explainability:** “[...] *Associated with the notion of explanation as an interface between humans and a decision maker [...].*” - **Use of these tools in a human interaction context.**



# The need of a formal definition

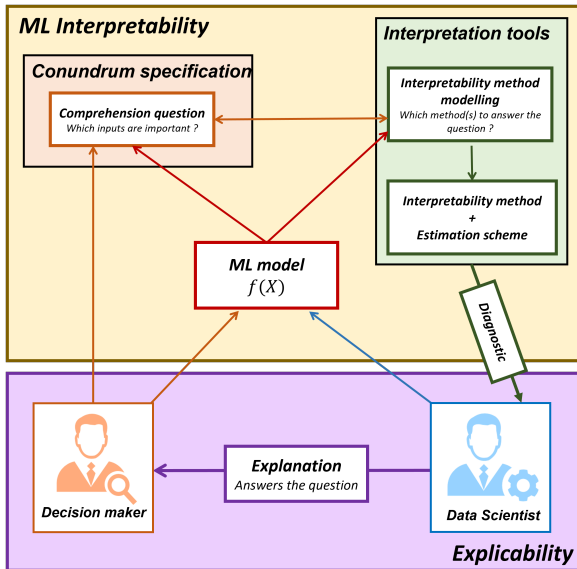
From the litterature (Barredo Arrieta et al. 2020):

- **Interpretability:** “[...] *The ability to explain or to provide the meaning in understandable terms to a human.*” - **Provide meaningful (mathematical) tools.**
- **Explainability:** “[...] *Associated with the notion of explanation as an interface between humans and a decision maker [...].*” - **Use of these tools in a human interaction context.**

Asher et al. (2022) provides a first step towards a formal model of explainability:

- 2-player explanation game between an **explainee** (decision maker) and an **explainer** (data scientist), to model the interaction:
  - Each player aim at solving a **conundrum** (a question on the black-box model).
  - The explainer provides **diagnostics** (interpretability method) at each turn.
  - The explainee either **accepts or rejects** the explanation (based, e.g., on the theoretical foundations of the interpretability method).

# Explanation game



# ML interpretability

In the rest of this presentation, ML interpretability can then be understood as:

- Development of interpretability methods relative to different conundrums.
- Theoretical understanding of their meaning and their validity.
- Development of efficient estimation schemes in order to compute relevant indices.

Example:

- **Conundrum:** Which feature is responsible for which part of a prediction ?
- **Interpretability method:** SHAP.
- **Estimation scheme:** KernelSHAP.
- **Diagnostic:** Compute the results given by KernelSHAP.
- **Explanation:** Interpret the results and provide an answer to the question.

# ML interpretability

In the rest of this presentation, ML interpretability can then be understood as:

- Development of interpretability methods relative to different conundrums.
- Theoretical understanding of their meaning and their validity.
- Development of efficient estimation schemes in order to compute relevant indices.

Example:

- **Conundrum:** Which feature is responsible for which part of a prediction ?
- **Interpretability method:** SHAP.
- **Estimation scheme:** KernelSHAP.
- **Diagnostic:** Compute the results given by KernelSHAP.
- **Explanation:** Interpret the results and provide an answer to the question.

**But is SHAP the only interpretability method that answers this question ? Is it the best ? What are its limits ?**

1. Explanation vs. Interpretation
2. Interpretability methods
3. Interpretability and cooperative game theory

# Meaningfulness of interpretability methods

In order to decide if an interpretability method is **meaningful**, i.e., if it should be accepted by the decision maker, its **theoretical properties must be understood**.

**Example:** What does SHAP quantify ?

Hence, one has to focus on the **theoretical quantities**, in order to provide **mathematical proofs** of the behavior of the interpretability method.

**Example:** When the features are independent, SHAP quantifies...

These properties allow to **accurately interpret the diagnostic**, and further enhances the impact of the **explanation**.

## Parenthesis: Statistical inference

**Goal:** Make a distinction between a **theoretical quantity** and an **estimator**.

One is interested in the average salary in France, denoted  $\mu$ . If the salary in France is assumed to be a random variable  $X$ , one would have that:

$$\mathbb{E}[X] = \mu.$$

This theoretical quantity could be measured exactly if **everyone in France would indicate its salary**. However, this is very costly.

In order to have an **educated guess** on  $\mu$ , i.e., a “good” approximation, one could randomly ask their salary to  $n$  persons living in France, resulting in  $x_1, \dots, x_n$  observations.

Then, an **estimator** of  $\mu$  would be the empirical mean:

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n x_i.$$

For instance, the theoretical values of **SHAP** can be estimated using several methods:  
**TreeSHAP, KernelSHAP...**

Here, we are interested in knowing if SHAP is a **meaningful theoretical quantity**, i.e., does it have theoretical properties that would make it the basis of a **good explanation** ?

In order to (partly) answer this question, one can start by taking an interest in the field of **Cooperative Game Theory**.



1. Explanation vs. Interpretation
2. Interpretability methods
3. Interpretability and cooperative game theory

# Cooperative game theory

In a nutshell, cooperative game theory can be summarized as “**the art of cutting a cake**”.



Given a **set of players**  $D = \{1, \dots, d\}$ , who produces a **quantity**  $v(D)$ , how can one allocate shares of  $v(D)$  among the  $d$  players ?

The “**cake cutting process**” is often described through **axioms** (i.e., desired properties), and results in an **allocation**.

Formally, a cooperative game is denoted  $(D, v)$  where  $D$  is a **set of players**, and  $v : \mathcal{P}(D) \rightarrow \mathbb{R}$  is a **value function**, mapping every possible subset of players to a real value.

# Cooperative game: illustration

Players

$D$



Coalitions

$\mathcal{P}_d$



Cost Function

$v$

Quantifies the  
value produced by  
a coalition

$$v(\text{blue}) = \text{eggs and flour}$$

$$v(\text{red}) = \text{chocolate and butter}$$

$$v(\text{yellow}) = \text{oven}$$

$$v(\text{blue}, \text{red}) = \text{cookies}$$

$$v(\text{red}, \text{yellow}) = \text{cake}$$

$$v(\text{yellow}, \text{blue}) = \text{bread}$$

$$v(\text{blue}, \text{red}, \text{yellow}) = \text{layer cake}$$

# Cooperative game: illustration

Players

$D$



Coalitions

$\mathcal{P}_d$



Cost Function

$v$

Quantifies the  
value produced by  
a coalition

$$v(\text{blue}) = \text{eggs and flour}$$

$$v(\text{red}) = \text{chocolate and butter}$$

$$v(\text{yellow}) = \text{oven}$$

$$v(\text{blue}, \text{red}) = \text{cookies}$$

$$v(\text{red}, \text{yellow}) = \text{cake}$$

$$v(\text{yellow}, \text{blue}) = \text{bread}$$

$$v(\text{blue}, \text{red}, \text{yellow}) = \text{cake with cherry}$$

How can we share the cake between the players ?

# Allocations

The cake-cutting process must respect one fundamental property:

- The **whole cake** and **nothing but the cake** must be shared (**Efficiency**).

For prediction decomposition (SHAP), we want **every feature contribution** (shares of cake) to **sum-up to the prediction itself** (the cake).

Example of efficient allocations:

- Give the whole cake to the yellow player.
- Divide the cake in three equal parts.
- Give half of the cake to blue, and one fourth to red and yellow.

# Allocations

The cake-cutting process must respect one fundamental property:

- The **whole cake** and **nothing but the cake** must be shared (**Efficiency**).

For prediction decomposition (SHAP), we want **every feature contribution** (shares of cake) to **sum-up to the prediction itself** (the cake).

Example of efficient allocations:

- Give the whole cake to the yellow player.
- Divide the cake in three equal parts.
- Give half of the cake to blue, and one fourth to red and yellow.

Can we define efficient allocations in a smarter way ?

**Random order model allocations** allow to easily define **efficient allocation**.

First, to every cooperative game  $(D, v)$  can be associated its **dual**  $(D, w)$  where  $\forall A \subseteq D$ :

$$w(A) = v(D) - v(D \setminus A).$$

**Random order model allocations** allow to easily define **efficient allocation**.

First, to every cooperative game  $(D, v)$  can be associated its **dual**  $(D, w)$  where  $\forall A \subseteq D$ :

$$w(A) = v(D) - v(D \setminus A).$$

Instead to quantifying the **value**  $v(A)$  produced by a coalition  $A$ , we rather focus on its **cost**  $w(A)$  (i.e., how much value did we lose by removing  $A$ ).  $w$  can be interpreted as the **bargaining power** of a coalition.



# The Weber set

**Random order model allocations** allow to easily define **efficient allocation**.

First, to every cooperative game  $(D, v)$  can be associated its **dual**  $(D, w)$  where  $\forall A \subseteq D$ :

$$w(A) = v(D) - v(D \setminus A).$$

Instead to quantifying the **value**  $v(A)$  produced by a coalition  $A$ , we rather focus on its **cost**  $w(A)$  (i.e., how much value did we lose by removing  $A$ ).  $w$  can be interpreted as the **bargaining power** of a coalition.


A random order allocation is computed as a **weighted average over the permutations of  $D$** : it considers every possible order each player can interact with the others.

# Backward-Forward and random order allocations

(a.)

Marginal value gain of adding players to  $\emptyset$ :  


$$v\left(C_{\pi(i)}(\pi)\right) - v\left(C_{\pi(i)-1}(\pi)\right)$$



(b.)

Marginal value cost of removing players from  $\{1, 2, 3\}$ :  

$$v\left(D \setminus C_{\pi(i)-1}(\pi)\right) - v\left(D \setminus C_{\pi(i)}(\pi)\right)$$



## Random order model allocations

For **any weighting scheme** (probability mass function) put on the permutations of  $D$  corresponds an **efficient allocation**.

## Random order model allocations

For **any weighting scheme** (probability mass function) put on the permutations of  $D$  corresponds an **efficient allocation**.

In this context, **to which weighting scheme corresponds the Shapley values ?**

## Random order model allocations

For **any weighting scheme** (probability mass function) put on the permutations of  $D$  corresponds an **efficient allocation**.

In this context, **to which weighting scheme corresponds the Shapley values ?**

They correspond to the **specific choice** of assigning the **same weight** to each permutation (i.e.,  $1/d!$ ).

# The Shapley values as a random order model allocation

The Shapley values of  $(D, v)$  and the ones of  $(D, w)$  are equal.

How can we understand the Shapley values as a random order model allocation:

- The order to which the players interact **does not matter**.
- The orders to which the players interact are **equally as likely** (this is a huge assumption).
- Allocating **costs** is the same as **allocating profits**.

They can be interpreted as

*“[...] an a priori assessment of the situation, based on either ignorance or disregard of the social organization of the players.” - L. S. Shapley (2016)*

**Unrelated but relevant parallel question:** Uniform prior in the Bayesian context.

## Proportional values

Defining an efficient allocation is **as easy as putting a weighting scheme on the permutations of  $D$** .

Are there other known allocations that might be relevant ?

# Proportional values

Defining an efficient allocation is **as easy as putting a weighting scheme on the permutations of  $D$** .

Are there other known allocations that might be relevant ?

Yes ! For instance, the **proportional values**. They correspond to a weighting scheme proportional to the **order of importance** of the permutations:

For a profit game  $(D, v)$ :

- 1 brings the most value ( $v(1)$  is bigger than the values of coalitions of 1 player).
- 3 added to 1 brings the most value ( $v(\{1, 3\})$  is bigger than the values of the coalitions of 2 players).
- 2 added to  $\{1, 3\}$  brings the most value (same thing for coalitions of 3 players).
- ...

Then the order  $(1, 3, 2, \dots)$  should have a **bigger weight** than the other orders. The same reasoning can be applied to the dual (cost) game  $(D, w)$ , resulting in a **different allocation**. 18/29



## Shapley values and proportional values in linear regression

These allocations have been used in statistics for a very long time (without the knowledge that they came from cooperative game theory).

In the context of linear regression, the question is to redistribute the  $R^2$  (the cake) among the players (the covariates).

# Shapley values and proportional values in linear regression

These allocations have been used in statistics for a very long time (without the knowledge that they came from cooperative game theory).

In the context of linear regression, the question is to redistribute the  $R^2$  (the cake) among the players (the covariates).

Let  $X_1, \dots, X_d$  be the features (vectors of  $n$  observations), and  $Y$  be the target. Denote  $R^2(A)$  be the  $R^2$  of the linear model:

$$Y = \beta_0 + \sum_{i \in A} \beta_i X_i + \epsilon.$$

## Shapley values and proportional values in linear regression

These allocations have been used in statistics for a very long time (without the knowledge that they came from cooperative game theory).

In the context of linear regression, the question is to redistribute the  $R^2$  (the cake) among the players (the covariates).

Let  $X_1, \dots, X_d$  be the features (vectors of  $n$  observations), and  $Y$  be the target. Denote  $R^2(A)$  be the  $R^2$  of the linear model:

$$Y = \beta_0 + \sum_{i \in A} \beta_i X_i + \epsilon.$$

In this case, the cake is  $v(D) = R^2(D)$ , the  $R^2$  of the linear regression of  $Y$  by every feature  $X_1, \dots, X_d$ .

And for any coalition of features  $A \subseteq D$ ,  $v(A) = R^2(A)$ , is the  $R^2$  of the **nested** linear regression of  $Y$  by only the coalition of features  $X_A$ .

$(D, R^2)$  forms a cooperative game !

# LMG and PMVD as importance measures

In order to have a **good** importance measure  $\phi_i, i = 1, \dots, d$ , in the context of linear regression (with correlated features), it must satisfy the four following criteria

- **Non-negativity:**  $\forall i \in D, \phi(i) \geq 0$  ;
- **Proper exclusion:** If  $\beta_i = 0$ , then  $\phi_i = 0$  ;
- **Proper inclusion:** If  $\beta_i \neq 0$ , then  $\phi_i > 0$  ;
- **Efficiency:**  $\sum_{i=1}^d \phi_i = R^2(D)$ .

Lindeman, Merenda, and Gold (1980) proposed indices called LMG, which are nothing more than **the Shapley values** of  $(D, R^2)$ . However they **violates the proper exclusion criterion**.

Feldman (2005) proposed indices called PMVD, which are nothing more than **the proportional values** of the dual of  $(D, R^2)$ . They **respect all four criterion**.

## Towards sensitivity analysis

But what if the model is non-linear, and the features are correlated ?

Then, whenever, for a non-linear model  $G$ , such that:

$$Y = G(X),$$

the cake is  $\mathbb{V}(Y) = \mathbb{V}(G(X))$  (notice that there is no  $\epsilon$ ). It leads to the **Sobol' cooperative game** where:

$$v(A) = \mathbb{V}(\mathbb{E}[G(X) \mid X_A]).$$

Its Shapley values are known as the **Shapley effects** (Owen 2014), and the proportional values of its dual are known as the **proportional marginal effects** (Herin et al. 2022).

Main distinction for correlated **exogenous** features:

- Their Shapley effects can be non-zero.
- Their proportional marginal effects are zero (this has been proven).


**Bottom line:** If the **conundrum** is about detecting exogenous inputs, the **Shapley effects may fail**, whereas the **proportional marginal effects do answer the question**.

## Interpreting the Shapley values: Harsanyi dividends

Another equivalent enlightening representation of the Shapley values can be done using **Harsanyi dividends** (Harsanyi 1963).

Let  $(D, v)$  be a cooperative game, and for any  $A \subseteq D$ , let the **Harsanyi dividend** of the coalition  $A$  be:

$$D_v(A) = \sum_{B \subseteq A} (-1)^{|A|-|B|} v(B).$$



The Harsanyi dividends can be interpreted as the **surplus (or shortfall)** that a coalition generates:

$$D_v(1) = v(1), \quad D_v(2) = v(2),$$


$$D_v(1, 2) = v(1, 2) - v(1) - v(2).$$

# Interpreting the Shapley values: Harsanyi dividends

The Shapley values are then defined as:

$$Sh_i = \sum_{A \subseteq D: i \in A} \frac{D_v(A)}{|A|},$$

or, in other words, each dividend of a coalition is **equally** redistributed between the players that composes it.



Quick example: Eve and John are two developers, Eve produces 10.000 lines of code, John produces 8.000 lines of code.


However, John really likes to play babyfoot, but Eve is a hard-worker.

When working together, they only produce 10.000 lines of code. This means that the dividend of their coalition is  $-8.000$ .

**Is it fair to attribute Eve  $-4.000$  lines of code, even if she did all the work ?**

## Interpreting the proportional values: Harsanyi dividends

The proportional values can also be interpreted in terms of Harsanyi dividends. It represents a **proportional** redistribution of the dividends of coalitions between the players that composes it.



Quick example: In this case, most of the –8.000 lines of code would be given to John, but Eve would also be receive a hit.

**Is it still fair ? Is it cooperative ?**



A lot of things are said in the literature about cooperative game theory, and the Shapley values, which are disputable:

- An allocation is a Shapley value.

A lot of things are said in the literature about cooperative game theory, and the Shapley values, which are disputable:

- An allocation is a Shapley value. **False.** Cooperative game theory is a rich field with many different solution concepts.

A lot of things are said in the literature about cooperative game theory, and the Shapley values, which are disputable:

- An allocation is a Shapley value. **False.** Cooperative game theory is a rich field with many different solution concepts.
- Cooperative game theory deals with cooperation.

A lot of things are said in the literature about cooperative game theory, and the Shapley values, which are disputable:

- An allocation is a Shapley value. **False.** Cooperative game theory is a rich field with many different solution concepts.
- Cooperative game theory deals with cooperation. **False.** It deals with ways coalitions of players can interact between each other. Example of a non-cooperative allocation: Proportional values.

A lot of things are said in the literature about cooperative game theory, and the Shapley values, which are disputable:

- An allocation is a Shapley value. **False.** Cooperative game theory is a rich field with many different solution concepts.
- Cooperative game theory deals with cooperation. **False.** It deals with ways coalitions of players can interact between each other. Example of a non-cooperative allocation: Proportional values.
- Shapley values are fair.

A lot of things are said in the literature about cooperative game theory, and the Shapley values, which are disputable:

- An allocation is a Shapley value. **False.** Cooperative game theory is a rich field with many different solution concepts.
- Cooperative game theory deals with cooperation. **False.** It deals with ways coalitions of players can interact between each other. Example of a non-cooperative allocation: Proportional values.
- Shapley values are fair. **False.** They are an egalitarian way of redistributing dividends.

A lot of things are said in the literature about cooperative game theory, and the Shapley values, which are disputable:

- An allocation is a Shapley value. **False.** Cooperative game theory is a rich field with many different solution concepts.
- Cooperative game theory deals with cooperation. **False.** It deals with ways coalitions of players can interact between each other. Example of a non-cooperative allocation: Proportional values.
- Shapley values are fair. **False.** They are an egalitarian way of redistributing dividends.
- Explanations based on Shapley values are always suitable.

A lot of things are said in the literature about cooperative game theory, and the Shapley values, which are disputable:

- An allocation is a Shapley value. **False.** Cooperative game theory is a rich field with many different solution concepts.
- Cooperative game theory deals with cooperation. **False.** It deals with ways coalitions of players can interact between each other. Example of a non-cooperative allocation: Proportional values.
- Shapley values are fair. **False.** They are an egalitarian way of redistributing dividends.
- Explanations based on Shapley values are always suitable. **False.** They answer a specific conundrum, which is yet to be fully understood (but we're working on it !).



# Conclusion

- Shapley values **can provide explanations**, but they are not **always suitable**.

# Conclusion

- Shapley values **can provide explanations**, but they are not **always suitable**.
- They rely on **egalitarian principles**, which is **not synonym to fairness**.

# Conclusion

- Shapley values **can provide explanations**, but they are not **always suitable**.
- They rely on **egalitarian principles**, which is **not synonym to fairness**.
- **Explainability** and **ML interpretability** is not restricted to Shapley values, allocations, or even **prediction decomposition**.

# Conclusion

- Shapley values **can provide explanations**, but they are not **always suitable**.
- They rely on **egalitarian principles**, which is **not synonym to fairness**.
- **Explainability** and **ML interpretability** is not restricted to Shapley values, allocations, or even **prediction decomposition**.
- It is not because an interpretation method gives a **good result**, that it is **always suitable**.

## To go further (and open questions)

- Consequences between a **player-centric** perspective (cooperative game theory) vs. a **model-centric** perspective (statistics/ML) ?

## To go further (and open questions)

- Consequences between a **player-centric** perspective (cooperative game theory) vs. a **model-centric** perspective (statistics/ML) ?
- Same cake as SHAP, but for another allocation ? Can we ensure nice properties ?

## To go further (and open questions)

- Consequences between a **player-centric** perspective (cooperative game theory) vs. a **model-centric** perspective (statistics/ML) ?
- Same cake as SHAP, but for another allocation ? Can we ensure nice properties ?
- Do the choice of **value function** matter in the diagnostic, and the subsequent explanation ?

## To go further (and open questions)

- Consequences between a **player-centric** perspective (cooperative game theory) vs. a **model-centric** perspective (statistics/ML) ?
- Same cake as SHAP, but for another allocation ? Can we ensure nice properties ?
- Do the choice of **value function** matter in the diagnostic, and the subsequent explanation ?
- Other forms of interpretability methods, for different types of conundrums ? (Robustness, stability, behavior guarantees...).



## To go further (and open questions)

- Consequences between a **player-centric** perspective (cooperative game theory) vs. a **model-centric** perspective (statistics/ML) ?
- Same cake as SHAP, but for another allocation ? Can we ensure nice properties ?
- Do the choice of **value function** matter in the diagnostic, and the subsequent explanation ?
- Other forms of interpretability methods, for different types of conundrums ? (Robustness, stability, behavior guarantees...).
- Come-up with, and improve **estimation strategies**.

# References i

- Asher, N., L. De Lara, S. Paul, and C. Russell. 2022. "Counterfactual Models for Fair and Adequate Explanations" [in en]. Number: 2 Publisher: Multidisciplinary Digital Publishing Institute, *Machine Learning and Knowledge Extraction* 4, no. 2 (June): 316–349. issn: 2504-4990, accessed November 28, 2022. <https://doi.org/10.3390/make4020014>. <https://www.mdpi.com/2504-4990/4/2/14>.
- Barredo Arrieta, A., N. Díaz-Rodríguez, J. Del Ser, A. Bennetot, S. Tabik, A. Barbado, S. Garcia, et al. 2020. "Explainable Artificial Intelligence (XAI): Concepts, taxonomies, opportunities and challenges toward responsible AI" [in en]. *Information Fusion* 58 (June): 82–115. issn: 1566-2535. <https://doi.org/10.1016/j.inffus.2019.12.012>. <https://www.sciencedirect.com/science/article/pii/S1566253519308103>.
- Feldman, B. E. 2005. "Relative Importance and Value" [in English]. *SSRN Electronic Journal*, issn: 1556-5068. <https://doi.org/10.2139/ssrn.2255827>. <http://www.ssrn.com/abstract=2255827>.
- Harsanyi, J. C. 1963. "A Simplified Bargaining Model for the n-Person Cooperative Game." *International Economic Review* 4 (2): 194–220. issn: 0020-6598.
- Herin, Margot, Marouane Idrissi, Vincent Chabridon, and Bertrand Iooss. 2022. "Proportional marginal effects for global sensitivity analysis." Working paper or preprint, October. <https://hal.archives-ouvertes.fr/hal-03825935>.
- Lindeman, R. H., P. F. Merenda, and R. Z. Gold. 1980. *Introduction to bivariate and multivariate analysis*. Glenview, IL: Scott Foresman / Company.
- Owen, A. B. 2014. "Sobol' Indices and Shapley Value" [in English]. *SIAM/ASA Journal on Uncertainty Quantification* 2, no. 1 (January): 245–251. issn: 2166-2525, accessed December 2, 2020. <https://doi.org/10.1137/130936233>.

Shapley, L. S. 2016. "17. A Value for n-Person Games." In *Contributions to the Theory of Games (AM-28), Volume II*, edited by Harold William Kuhn and Albert William Tucker, 307–318. Princeton University Press.  
<https://doi.org/doi:10.1515/9781400881970-018>.

**THANK YOU FOR YOUR ATTENTION!**

**ANY QUESTIONS?**